

Transport Layer

TCP/IP

Invalid Date

Table of contents

1	Transport Layer	3
1.1		3
1.2	Internet Transport Protocol	3
2	Multiplexing & Demultiplexing	3
2.1		3
2.2	UDP (Connectionless) Demux	4
2.3	TCP (Connection-oriented) Demux	4
3	UDP (User Datagram Protocol)	4
3.1	Segment	4
3.2	UDP	4
3.3	UDP Checksum	5
3.3.1		5
3.3.2	1 (One's Complement)	5
3.3.3	(16-bit)	5
3.3.4		6
4	Principles of Reliable Data Transfer (rdt)	6
4.1		6
4.2	rdt 2.0 —	6
4.3	rdt 3.0 — Stop-and-Wait	6
4.3.1	rdt 3.0 4	7
5	Pipelined Protocols	7
5.1		7
5.2	Go-Back-N (GBN) vs Selective Repeat (SR)	7
6	Window Size —	8
6.1		8
6.2	GBN	8
6.3	SR	8

6.4	SR Dilemma — $N/2$?	8
6.4.1	Seq#	8
6.4.2		9
7	TCP	9
7.1	Segment	9
7.2	TCP Sequence Number & ACK	10
7.2.1	TCP “byte stream”	10
7.2.2	Telnet (p.54)	10
7.3	RTT	11
7.3.1	EWMA (Exponentially Weighted Moving Average)	11
7.3.2	(DevRTT)	11
7.3.3	Timeout Interval	11
7.3.4	Retransmission RTT ?	11
7.4	TCP Reliable Data Transfer	12
7.4.1		12
7.4.2	ACK (RFC 1122, 2581)	12
7.4.3	Cumulative ACK (p.62)	12
7.4.4	Fast Retransmit	12
7.5	TCP Flow Control	13
7.5.1		13
7.5.2		13
7.5.3	rwnd = 0	13
7.5.4	Flow Control vs Congestion Control	13
7.6	TCP Connection Management	14
7.6.1	3-way Handshake	14
7.6.2	Connection Teardown	14
8	Principles of Congestion Control	14
8.1	(Congestion) ?	14
8.2	(3)	14
8.2.1	Scenario 2 : Goodput vs Throughput	14
8.2.2	Scenario 3: Upstream	15
9	TCP Congestion Control	15
9.1	AIMD (Additive Increase Multiplicative Decrease)	15
9.2	Slow Start	15
9.3	TCP Tahoe vs TCP Reno	16
9.3.1		16
9.3.2	TCP Tahoe (1988)	16
9.3.3	TCP Reno (1990)	16
9.3.4	cwnd	17
9.3.5	TCP	17
9.4	TCP Throughput	17
9.5	TCP Fairness	18
9.6	ECN (Explicit Congestion Notification)	18
10		18

1 Transport Layer

1.1

Transport layer **application process** (logical communication) .

Network layer	host	host
Transport layer	process	process

- : App segment → Network layer
- : Segment → App layer

i : Household Analogy

Ann 12 Bill 12 :

- Houses = **hosts**
- Kids = **processes**
- Letters = **app messages**
- Ann/Bill () = **transport protocol (demux)**
- = **network-layer protocol**

1.2 Internet Transport Protocol

	TCP	UDP
	Connection-oriented	Connectionless
	Flow control	
	Congestion control	
Header	20 bytes	8 bytes
	HTTP, FTP, SMTP	DNS, streaming, SNMP

2 Multiplexing & Demultiplexing

2.1

- **Multiplexing** (): transport header network layer
- **Demultiplexing** (): header segment

2.2 UDP (Connectionless) Demux

- **2-tuple:** (dest IP, dest port)
- (src IP, src port) dest port →

2.3 TCP (Connection-oriented) Demux

- **4-tuple:** (src IP, src port, dest IP, dest port)
- dest port 80 src demux



TCP 4-tuple , (80) .

3 UDP (User Datagram Protocol)

3.1 Segment

source port #	dest port #	16 + 16 bits
length	checksum	16 + 16 bits
application data		

- **length:** UDP segment (header + data)

3.2 UDP

- No connection establishment →
- No connection state →
- Small header (8 bytes vs TCP 20 bytes)
- No congestion control →

3.3 UDP Checksum

3.3.1

“segment ” , Pseudo Header + UDP Header + UDP Data .

Pseudo Header (IP header):

Source IP Address (32 bits)

Destination IP Address (32 bits)

zero (8b) Protocol=17 (8b)

UDP Length (16 bits)

i Pseudo Header ?

UDP header IP . Pseudo header **IP** (wrong delivery) .

3.3.2 1 (One's Complement)

1 :

0110 1010 → 1001 0101 (1)

$$: x + \sim x = \underbrace{1111 \dots 1}_{\text{all ones}}$$

3.3.3 (16-bit)

1111 0011 0011 0011 0 ← (17 , carry)
+ 1110 1010 1010 1010 1

1101 1101 1101 1101 1 ← carry out
+ 1 ← wrap around (carry)

1101 1101 1101 1110 0 ← sum

Checksum = 1 :

0010 0010 0010 0001 1 ← checksum

3.3.4

: 16-bit word + checksum

$$\text{sum} + \text{checksum} = \underbrace{1111 \dots 1}_{\text{all ones}} \Rightarrow$$

$$\text{sum} + \text{checksum} \neq \underbrace{1111 \dots 1}_{\text{all ones}} \Rightarrow$$



. (best-effort)

4 Principles of Reliable Data Transfer (rdt)

4.1

rdt 1.0	-	(trivial)
rdt 2.0	checksum, ACK/NAK	
rdt 2.1	Sequence number (0,1)	ACK/NAK
rdt 2.2	NAK , duplicate ACK	NAK-free
rdt 3.0	Countdown timer	

4.2 rdt 2.0 —

ACK/NAK sender .

- → duplicate
- → **Sequence number** (rdt 2.1)

4.3 rdt 3.0 — Stop-and-Wait

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R}$$

: 1 Gbps link, RTT = 30ms, L = 8000 bits

$$U_{\text{sender}} = \frac{0.008}{30.008} \approx 0.00027 \quad (0.027\%!)$$

→

4.3.1 rdt 3.0 4

(a)	-	
(b) Packet loss		timeout →
(c) ACK loss	ACK	timeout → , receiver duplicate re-ACK
(d) Premature timeout	ACK	seq# duplicate

5 Pipelined Protocols

5.1

3 :

$$U_{sender} = \frac{3 \cdot L/R}{RTT + L/R}$$

→ 3

5.2 Go-Back-N (GBN) vs Selective Repeat (SR)

	GBN	SR
Sender window	N in-flight	N in-flight
Receiver window	1	N
ACK	Cumulative ACK	ACK
Out-of-order	Discard () window unacked pkt	Buffer pkt
Timer	oldest unacked pkt	pkt

6 Window Size —

6.1

window size :

$$W_{sender} + W_{receiver} \leq \text{Seq\# space (N)}$$

Sender seq# Receiver “ ” seq# .

6.2 GBN

GBN receiver window = 1 :

$$W + 1 \leq N \Rightarrow W \leq N - 1 = 2^k - 1$$

- k : seq# , $N = 2^k$: seq#

6.3 SR

SR receiver window = sender window = **W** :

$$W + W \leq N \Rightarrow W \leq \frac{N}{2}$$

!

GBN ($W \leq 2^k - 1$) SR ($W \leq 2^k/2$) . GBN receiver window 1 .
() , receiver window sender window .

6.4 SR Dilemma — N/2 ?

6.4.1 Seq#

Seq# (circular) . (N=4: 0,1,2,3)

W = 3 ($> N/2 = 2$) $\rightarrow$


Step 1: Sender [0,1,2] $\rightarrow$ pkt 0,1,2

Step 2: ACK 0,1,2

Receiver window : [3,0,1] $\leftarrow$

Step 3: Sender timeout $\rightarrow$ pkt 0

: seq# 0 receiver window [3,0,1] !
Receiver: " "




Receiver window = [3, 0, 1] ← 0 (!)
 Sender window = [0, 1, 2]

$W = 2 (= N/2) \rightarrow$

Step 1: Sender [0,1] → pkt 0,1
 Step 2: ACK 0,1
 Receiver window : [2,3] ←
 Step 3: Sender timeout → pkt 0

: seq# 0 receiver window [2,3] !
 Receiver: "old data → "



Receiver window = [2, 3] ← 0 (!)
 Sender window = [0, 1]

6.4.2

window N .

$$2W > N \Rightarrow \text{ } > \Rightarrow$$

$$2W \leq N \Rightarrow \text{non-overlapping}$$

7 TCP

7.1 Segment

source port dest port 16 + 16 bits

sequence number	32 bits	
acknowledgement number	32 bits	
hlen	flags	receive window
checksum	urg	data pointer
options (variable)		
application data		

: URG, ACK, PSH, RST, SYN, FIN

7.2 TCP Sequence Number & ACK

7.2.1 TCP “byte stream”

TCP .

- **Seq#**: segment
- **ACK#**: (= + 1)
- **Cumulative ACK**: “ ”

7.2.2 Telnet (p.54)

: A seq# = 42, B seq# = 79

A [Seq=42, ACK=79, data='C'] B

Seq=42: " 42 "
 ACK=79: "B 79 "
 data='C': 1

A [Seq=79, ACK=43, data='C'] B

Seq=79: " 79 "
 ACK=43: "A 42 (1)
 → 43 "

A [Seq=43, ACK=80] B

ACK=80: "B 79 (1)
 → 80 "

💡 ACK = Seq + data

ACK# "X" = "X-1" . Seq=42, 1 → ACK=43 (42+1).

7.3 RTT

7.3.1 EWMA (Exponentially Weighted Moving Average)

$$\text{EstimatedRTT} = (1 - \alpha) \cdot \text{EstimatedRTT} + \alpha \cdot \text{SampleRTT}$$

$$\alpha = 0.125 \text{ (typical)}$$

→ .

7.3.2 (DevRTT)

$$\text{DevRTT} = (1 - \beta) \cdot \text{DevRTT} + \beta \cdot |\text{SampleRTT} - \text{EstimatedRTT}|$$

$$\beta = 0.25 \text{ (typical)}$$

7.3.3 Timeout Interval

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 \cdot \text{DevRTT}$$

i

mean + 4σ , 99.99% . safety margin .

7.3.4 Retransmission RTT ?

Ambiguity (Karn's Algorithm):

t=0: pkt (seq=100)

t=5: timeout →

t=8: ACK(100)

SampleRTT = ?

: 8 - 0 = 8ms ()

: 8 - 5 = 3ms ()

→ ACK !

RTT : - → Timeout → - → Timeout →

7.4 TCP Reliable Data Transfer

7.4.1

1. **Timeout:** unacked segment
2. **3x Duplicate ACK:** Fast Retransmit

7.4.2 ACK (RFC 1122, 2581)

segment, pending ACK	500ms ACK (Delayed ACK)
segment, pending ACK	cumulative ACK
Out-of-order segment (Gap)	duplicate ACK
Gap segment	ACK

7.4.3 Cumulative ACK (p.62)

Host A

Host B

Seq=92 (8 bytes)

Seq=100 (20 bytes)

← ACK=100

(B 92, 100 buffer)

timeout !

Seq=92

ACK=120

ACK=120 ? B Seq=92(8 bytes) Seq=100(20 bytes) buffer : - 92~99 (8 bytes) +
100~119 (20 bytes) = **119** - → ACK=120: “120 (119)”

Cumulative ACK : ACK=100 , ACK=120 100~119 .

7.4.4 Fast Retransmit

3 duplicate ACK → timeout

A [Seq=92] B
A [Seq=100] X ()
A [Seq=120] B
A [Seq=140] B

B ACK=100 (dup)

B ACK=100 (dup)

B ACK=100 (dup) ← 3 !

A: timeout Seq=100

7.5 TCP Flow Control

7.5.1

buffer overflow (cf. Congestion control)

7.5.2

$$rwnd = RcvBuffer - [LastByteRcvd - LastByteRead]$$

- **RcvBuffer:** (4096 bytes)
- **LastByteRcvd:**
- **LastByteRead:** Application

Receiver TCP segment header **rwnd** .

Sender :

$$LastByteSent - LastByteAcked \leq rwnd$$

7.5.3 rwnd = 0

Receiver rwnd=0 sender . buffer sender (trigger).

: sender **1 byte probe segment** . Receiver ACK rwnd .

7.5.4 Flow Control vs Congestion Control

Flow Control	Congestion Control
buffer	
rwnd (receiver → sender)	packet loss, delay
rwnd sender	cwnd sender
min(rwnd, cwnd)	

7.6 TCP Connection Management

7.6.1 3-way Handshake

Client	Server
[SYN, Seq=x]	
[SYNACK, Seq=y, ACK=x+1]	
[ACK, ACK=y+1]	

Client state: CLOSED → SYN_SENT → ESTAB

Server state: LISTEN → SYN_RCVD → ESTAB

2-way handshake : , , half-open connection .

7.6.2 Connection Teardown

Client	Server
[FIN]	Client: FIN_WAIT_1 → FIN_WAIT_2
[ACK]	Server: CLOSE_WAIT
[FIN]	Server: LAST_ACK
[ACK]	Client: TIMED_WAIT (2×MSL) → CLOSED
	Server: CLOSED

TIMED_WAIT: $2 \times \text{MSL}$ (Maximum Segment Lifetime) . duplicate segment .

8 Principles of Congestion Control

8.1 (Congestion) ?

“ source ”

Flow control : - Flow control: buffer - Congestion control:

8.2 (3)

8.2.1 Scenario 2 : Goodput vs Throughput

Perfect knowledge ():

$\text{lin}' = R/2$:
 $=$ + ()
 $\text{drop} =$
 $\rightarrow R/2$
 $\text{goodput} = R/2$ (asymptotic)

💡 p.83

“sending at $R/2$ goodput $R/2$?”
 () . $R/2$.
premature timeout bandwidth goodput $< R/2$.

Realistic (timeout):

$$\text{goodput} < \frac{R}{2}$$

8.2.2 Scenario 3: Upstream

drop , capacity .

9 TCP Congestion Control

9.1 AIMD (Additive Increase Multiplicative Decrease)

- **Additive Increase:** loss RTT cwnd += 1 MSS
- **Multiplicative Decrease:** loss cwnd = cwnd / 2

$$\text{sending rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

9.2 Slow Start

- cwnd = 1 MSS
- RTT 2 (ACK +1 MSS →)
- ssthresh linear(CA)

cwnd

$\leftarrow \text{ssthresh}$

(linear)

()

time

(Slow Start)

9.3 TCP Tahoe vs TCP Reno

9.3.1

loss ?

Timeout	3 Dup ACK
	, 3 →

9.3.2 TCP Tahoe (1988)

loss :

Timeout : cwnd = 1 MSS, ssthresh = cwnd_before/2, Slow Start

3 Dup ACK : cwnd = 1 MSS, ssthresh = cwnd_before/2, Slow Start

9.3.3 TCP Reno (1990)

3 Dup ACK **Fast Recovery** :

Timeout :

→ cwnd = 1 MSS

→ ssthresh = cwnd / 2

→ Slow Start (Tahoe)

3 Dup ACK (Fast Recovery):

→ ssthresh = cwnd / 2

→ cwnd = ssthresh + 3 MSS $\leftarrow$ 3 dup ACK = 3 pkt

→ segment

→ dup ACK : cwnd += 1 MSS (window inflation)


→ New ACK : cwnd = ssthresh, CA

i Window Inflation

Fast Recovery :
 dup ACK dup ACK 1 sender 1

9.3.4 cwnd

cwnd



Tahoe: loss → (1 MSS)

Reno: 3dup → ssthresh ,

9.3.5 TCP

Tahoe	1988	congestion control, loss = cwnd=1
Reno	1990	Fast Recovery (3 dup ACK)
NewReno	1999	Reno partial ACK
CUBIC	2008	Linux , cubic cwnd
BBR	2016	bandwidth & RTT (Google)

9.4 TCP Throughput

Window size W :

$$\text{avg throughput} = \frac{3}{4} \cdot \frac{W}{\text{RTT}}$$

Loss L (Mathis 1997):

$$\text{TCP throughput} = \frac{1.22 \cdot \text{MSS}}{\text{RTT} \cdot \sqrt{L}}$$

: MSS=1500B, RTT=100ms, 10 Gbps:

$$L \leq 2 \times 10^{-10} \quad (\text{loss rate})$$

9.5 TCP Fairness

K TCP bandwidth $R \rightarrow R/K$

AIMD : throughput “equal bandwidth” zigzag (additive increase = 1 , multiplicative decrease =)

9.6 ECN (Explicit Congestion Notification)

- IP ToS field 2
- TCP ACK **ECE bit** $\rightarrow$
-

10

! TCP

TCP Tahoe/Reno **loss = congestion** .
 : - $\rightarrow$ (!) - $\rightarrow$ - (fading) $\rightarrow$ loss
 TCP congestion $\rightarrow$ cwnd $\rightarrow$
 BBR, CUBIC, TCP (: TCP-Westwood, TCP-Veno) .

11

$U_{sender} = \frac{L/R}{RTT+L/R}$	Stop-and-wait sender utilization
$W_{GBN} \leq 2^k - 1$	GBN window size
$W_{SR} \leq \frac{2^k}{2}$	SR window size
$EstRTT = (1 - \alpha)EstRTT + \alpha \cdot SampleRTT$	EWMA RTT ($\alpha = 0.125$)
$DevRTT =$	RTT ($\beta = 0.25$)
$(1 - \beta)DevRTT + \beta \cdot \ SampleRTT - EstRTT\ $	
$TimeoutInterval = EstRTT + 4 \cdot DevRTT$	TCP timeout interval
$rwnd =$	Flow control
$RcvBuffer - [LastByteRcvd - LastByteRead]$	
$avg\ throughput = \frac{3W}{4 \cdot RTT}$	TCP
$TCP\ throughput = \frac{1.22 \cdot MSS}{RTT \sqrt{L}}$	Loss (Mathis)